



# HADOOP DAEMONS

- Hadoop Terminology
- Hadoop 1.0 Daemons
- Hadoop 2.0 Daemons
- References

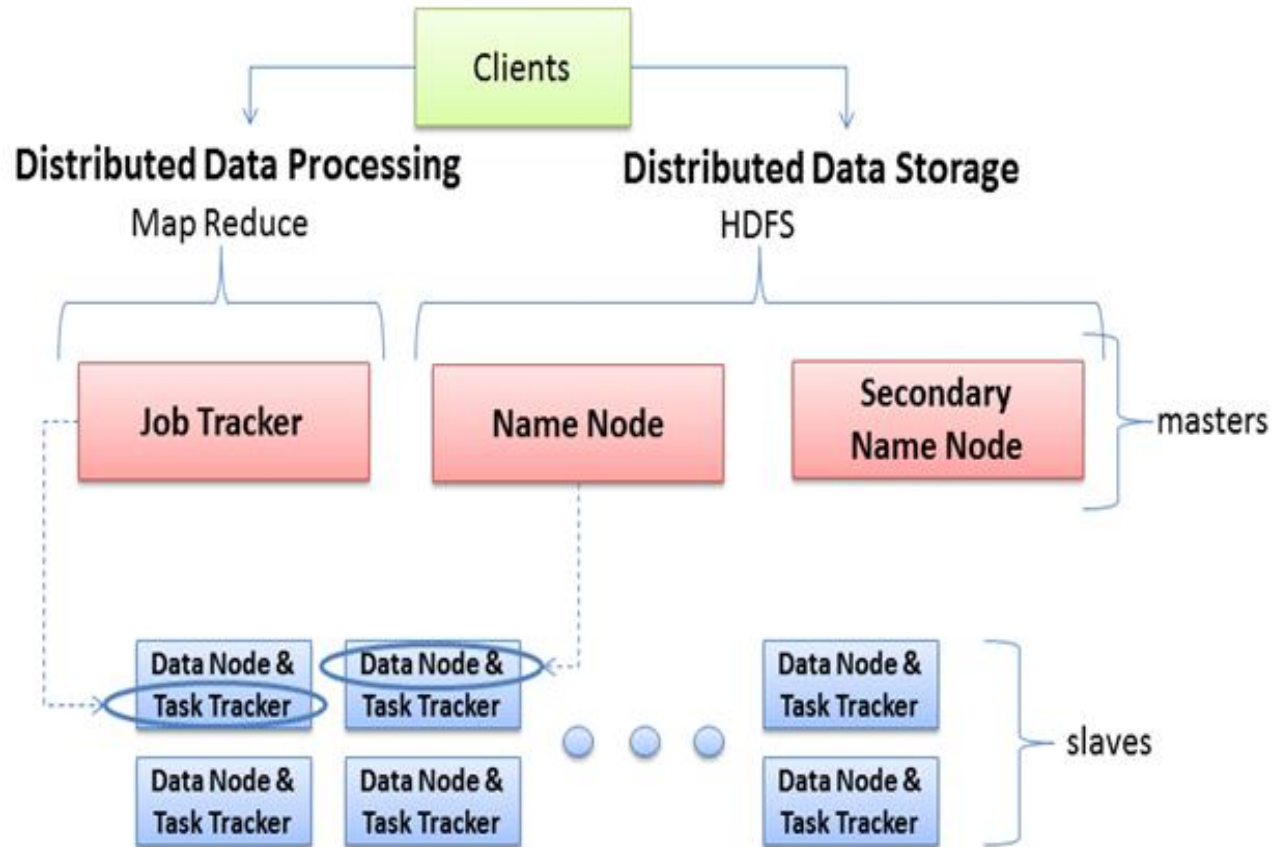
# Hadoop Terminology

Terminology	Definition
<b>Job</b>	A full program – an execution of a Mapper and Reducer across data set
<b>Task</b>	An execution of a mapper or reducer on a slice of data.
<b>Task Attempt</b>	Particular instance of an attempt to execute a task on a machine
<b>Daemon</b>	Computer program that runs as a background process rather than being under the direct control of an interactive user ( like Services in Windows )

# Hadoop 1.0 Daemons

## Daemons in Hadoop 1.0

- NameNode
- Secondary NameNode
- DataNode
- JobTracker
- TaskTracker



# NameNode

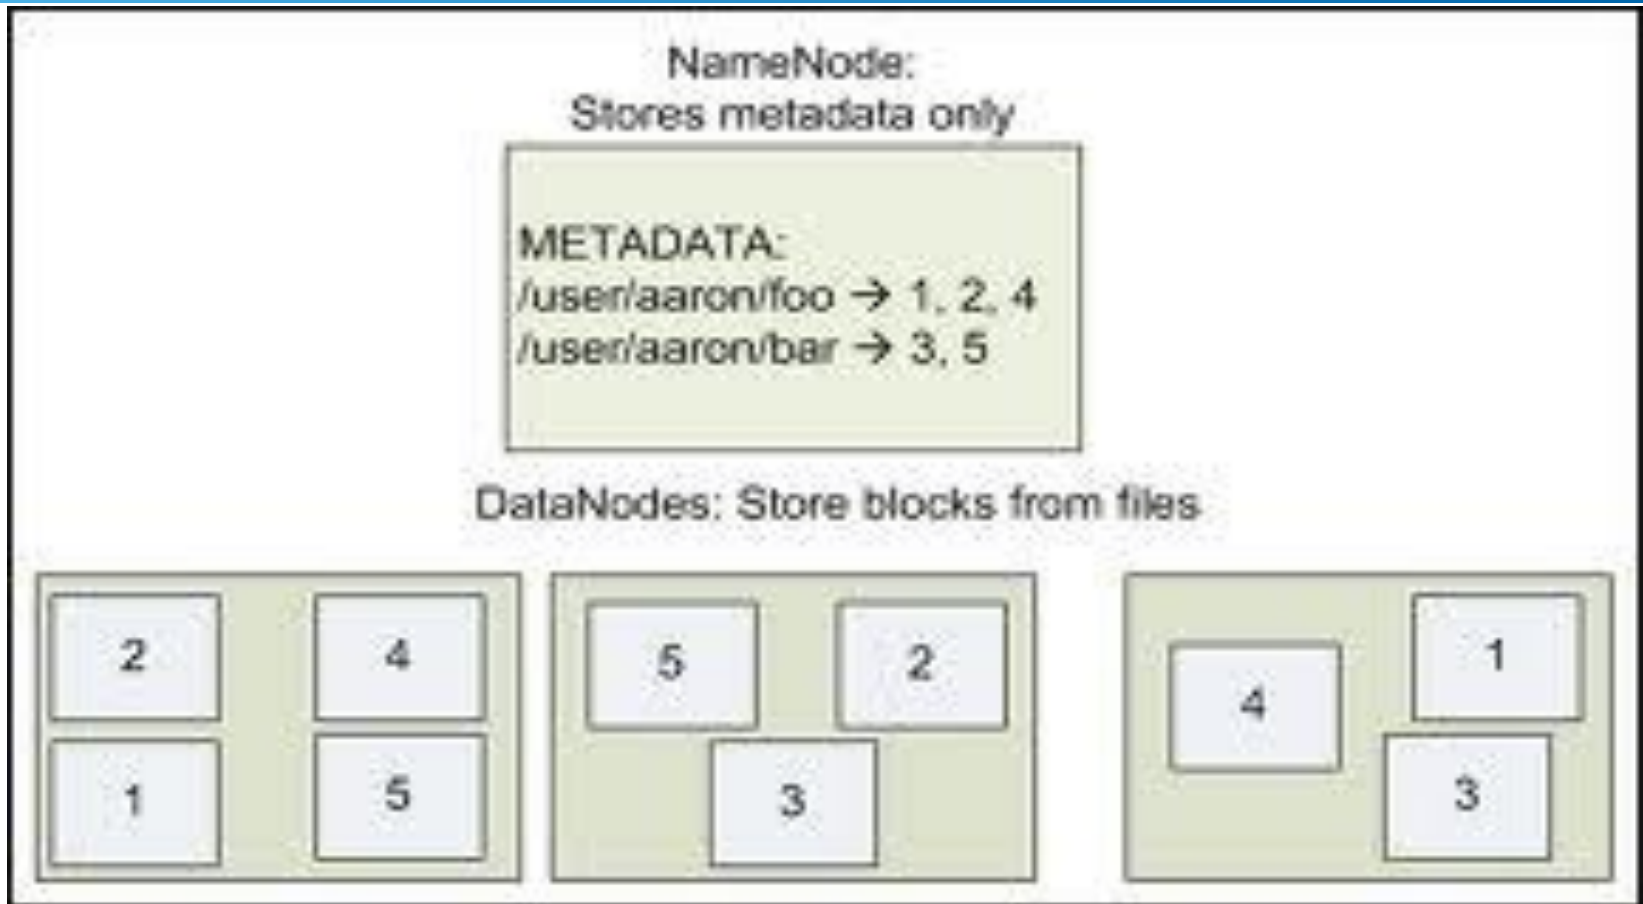
- ✓ Centerpiece of HDFS.
- ✓ <http://localhost:50070/>
  - web UI for HDFS name node(s)
- ✓ Maintains directory tree files in file system.
- ✓ Tracks file data in cluster.
  - Does not store the data of these files itself.
- ✓ Client application talk to NameNode to locate a file.
- ✓ Responds to successful client requests.
- ✓ Returns list of relevant DataNodes where data lives.
- ✓ Use heartbeats to detect DataNode failure.
- ✓ Chooses new DataNodes for new replicas.
- ✓ Single point of failure for the HDFS Cluster.

## Secondary NameNode

- ✓ Assistant daemon for monitoring the state of HDFS.
- ✓ Each cluster has one SNN residing on its own machine.
- ✓ No other DataNode or TaskTracker runs on SNN.
- ✓ Doesn't receive/record any real-time changes to HDFS.
- ✓ Allow NameNode to preserve file system's metadata to act as a backup.
- ✓ Communicates with NameNode.
- ✓ Take snapshots of metadata at intervals defined by the cluster configuration.
- ✓ Snapshots minimize the downtime and loss of data.

- ✓ Individual machines in a cluster are referred to as DataNodes.
- ✓ In HDFS files are broken into blocks of fixed size.
- ✓ These blocks are stored in DataNodes
- ✓ DataNode communicate with other DataNodes to replicate data blocks for redundancy.
- ✓ DataNode connects to NameNode to establish service.
- ✓ Responds to requests from the NameNode.
- ✓ DataNodes send heartbeat to the NameNode.

# How Data is stored in HDFS ?



DataNodes holding multiple file blocks with a replication factor of 2  
NameNode maps the filenames into the block ids

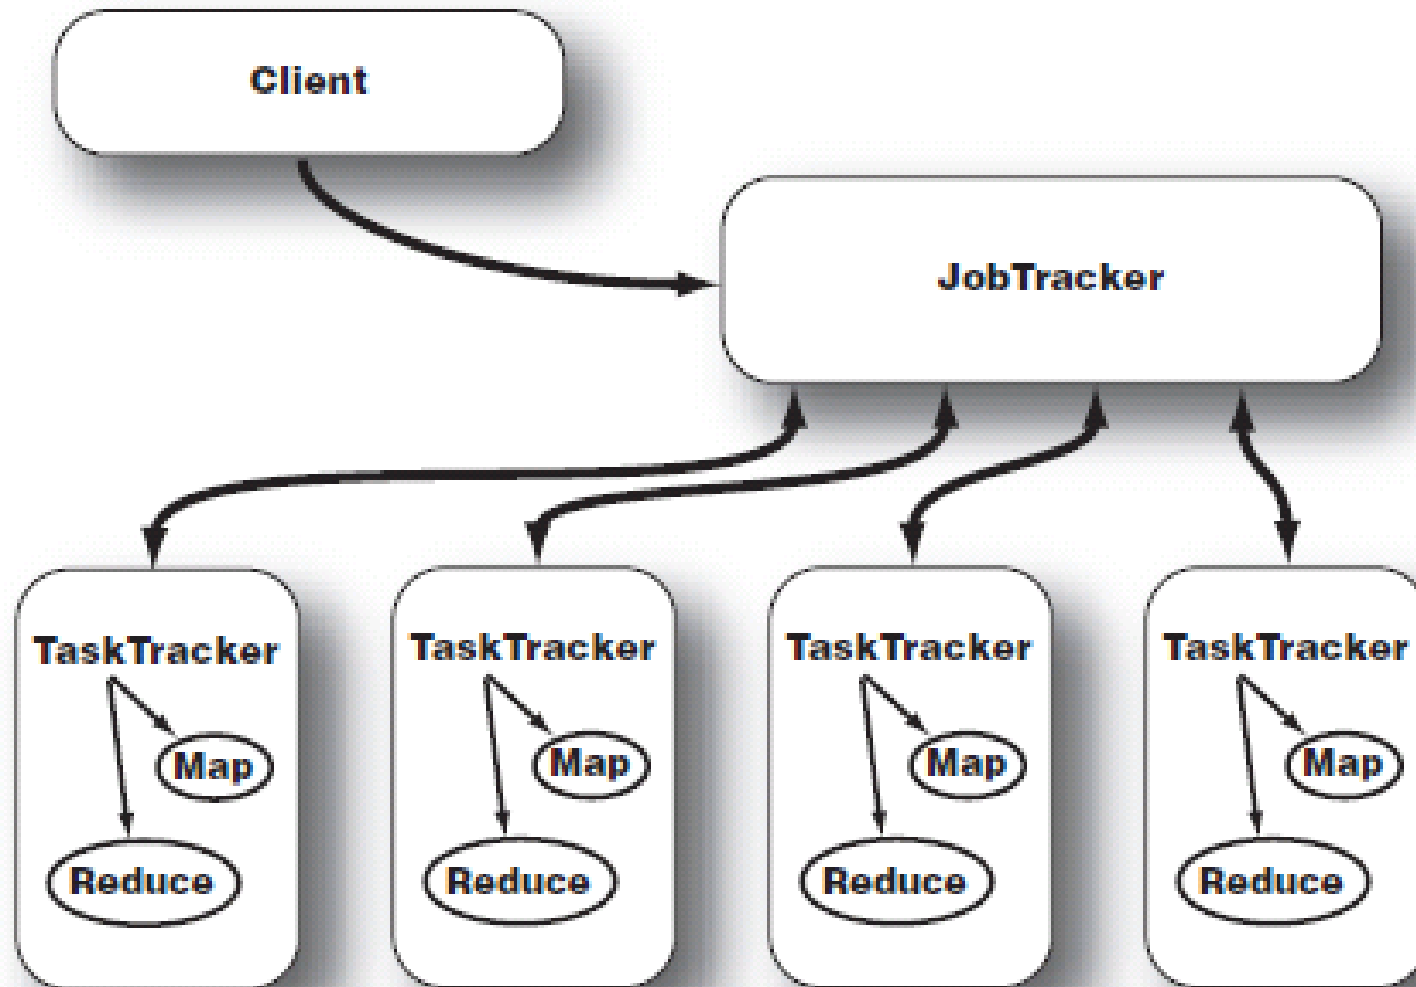
Default Replication factor is 3



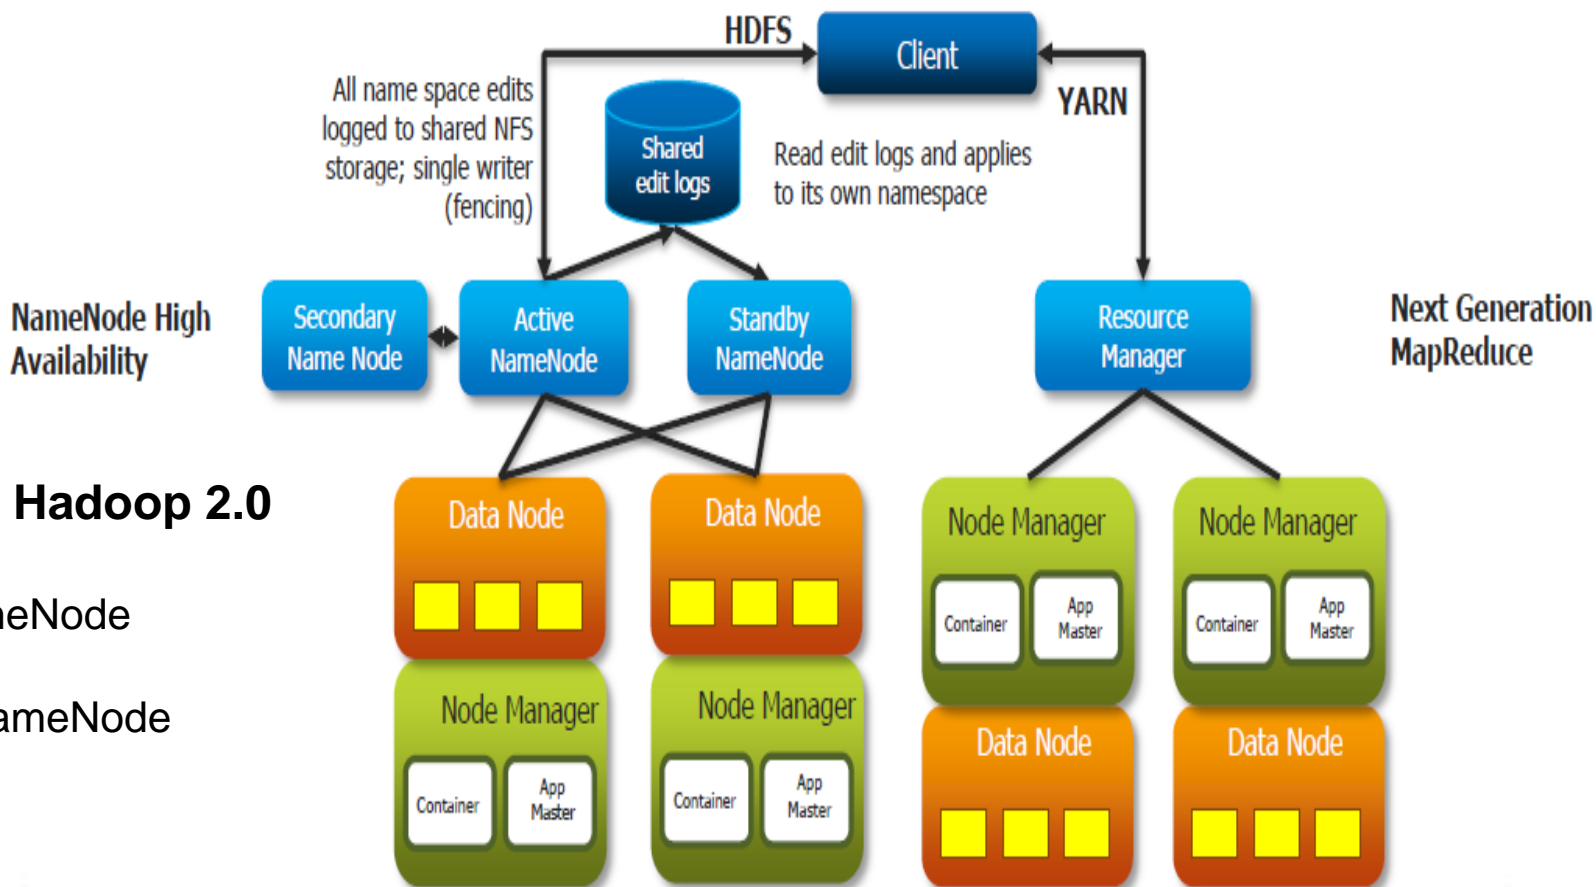
- ✓ Interface between application and Hadoop
- ✓ Run on a server as master node of the cluster.
- ✓ Determines the execution plan for file processing.
- ✓ Assigns nodes to different tasks and monitors them.
- ✓ If a task fails, re-launches the task on a different node.
- ✓ Only one JobTracker per cluster.
- ✓ <http://localhost:50030/>
  - web UI for MapReduce JobTracker(s)

- ✓ Manage the execution of individual tasks on each slave node.
- ✓ Single TaskTracker per slave node.
- ✓ Can handle many map or reduce tasks in parallel.
- ✓ Constantly communicates with JobTracker.
- ✓ In case of failure of heartbeat reception from tasktracker
  - JobTracker assumes TaskTracker has crashed.
  - Resubmit corresponding tasks to another.
- ✓ <http://localhost:50060/>
  - web UI for task tracker(s)

# JobTracker and TaskTracker Interaction



# Hadoop 2.0 Daemons



## Daemons in Hadoop 2.0

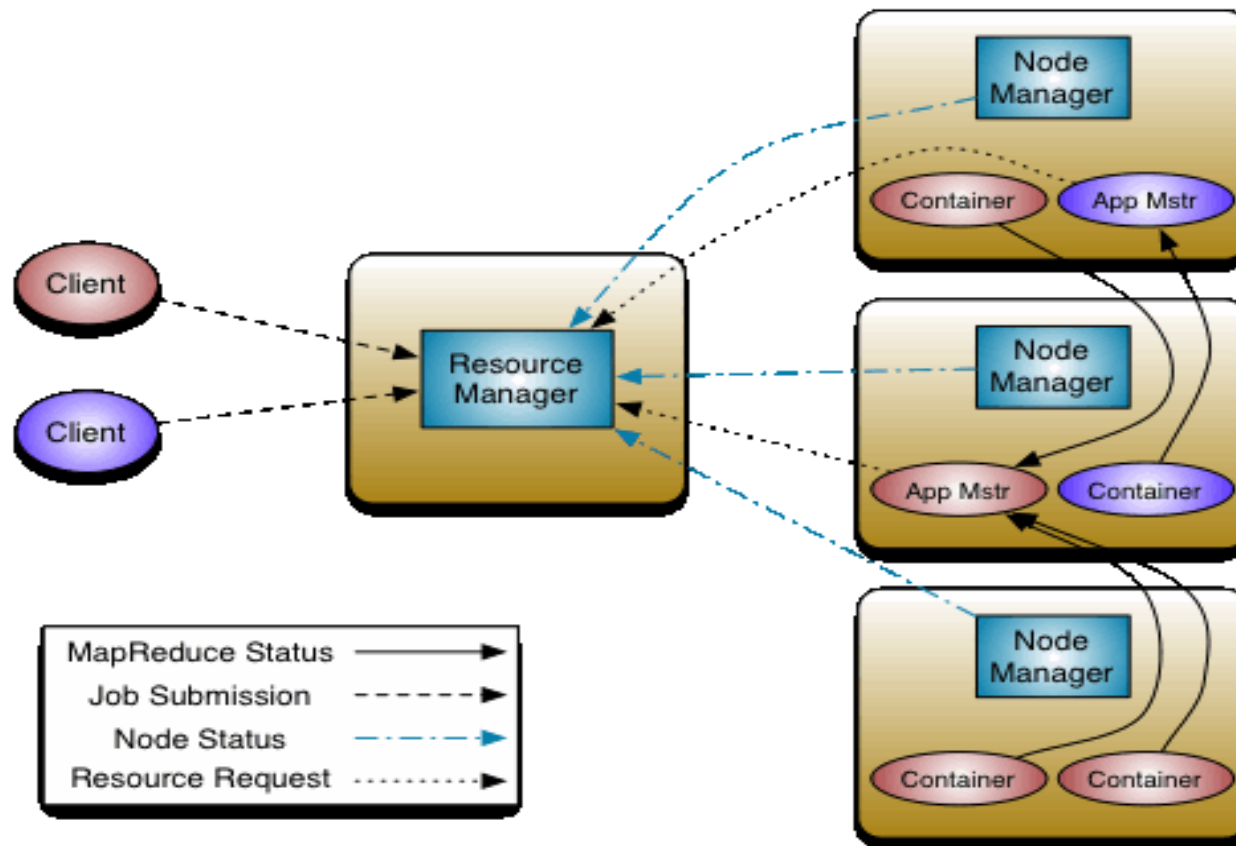
- Active NameNode
- Standby NameNode
- DataNode
- Resource Manager
- Node Manager

# Hadoop 2.0 Daemons

- ✓ NameNode was a single point of failure (SPOF) in an Hadoop 1.0 cluster
- ✓ Each cluster had a single NameNode, and if that machine or process became unavailable, the cluster as a whole would be unavailable until the NameNode was either restarted or brought up on a separate machine
- ✓ In Hadoop 2.0, we have two redundant NameNodes in the same cluster in an Active/Passive configuration with a hot standby
- ✓ Exactly one of the NameNodes is in an Active state, and the other is in a Standby state
- ✓ The Active NameNode is responsible for all client operations in the cluster, while the Standby is simply acting as a slave, maintaining enough state to provide a fast failover if necessary

- ✓ An Application is either a single job in the classical sense of Map-Reduce jobs or a DAG of jobs
- ✓ The ResourceManager and NodeManager (NM), form the data-computation framework
- ✓ The ResourceManager is the ultimate authority that arbitrates resources among all the applications in the system
- ✓ The NodeManager is the per-machine framework agent who is responsible for containers, monitoring their resource usage (CPU, memory, disk, network) and reporting the same to the ResourceManager/Scheduler
- ✓ The ApplicationMaster is a framework specific library and is tasked with negotiating resources from the ResourceManager and working with the NodeManager(s) to execute and monitor the tasks
- ✓ The MapReduce job history server allow the user to get status on finished applications


# ResourceManager & NodeManager Interaction



- Hadoop Wiki.
- Yahoo Hadoop Tutorials.
- Introduction to HDFS, Developer Works, IBM.
- Hadoop In Action



Thank You



**TATA**

Promise what we deliver.  
Deliver what we promise. That's  
**certainty**

Critical situations. Ruthless competition. Unforgiving customers. Thankfully you can be absolutely sure of your IT solutions with Tata Consultancy Services (TCS). As one of the world's fastest growing technology and business solutions providers, TCS has built a reputation of delivery excellence based on world class IT solutions that are on time, within budget and consistently deliver superior quality. So it comes as no surprise that we pioneered the concept of the Global Network Delivery Model, Developed Innovation Labs and Solution Accelerators. Achieving a level of delivery excellence that provides greater value to our customers and is the industry benchmark. Enabling our clients to experience certainty.

**TATA CONSULTANCY SERVICES**  
Experience certainty.  
IT Services • Business Solutions • Outsourcing

To learn how your business can experience certainty, visit [www.tcs.com](http://www.tcs.com)